

Course = BIL 470 / HOMEWORK 2

Name = MERT CAN GÖNEN

ID = 181101039

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn import metrics
from LR import LinearRegressionModel
```

Exploratory Data Analysis (EDA)

Read Dataset

```
In [2]: data = pd.read_csv("data.csv");
```

```
In [3]: display(data);
```

	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3
...
495	Female	150	153	5
496	Female	184	121	4
497	Female	141	136	5
498	Male	150	95	5
499	Male	173	131	5

500 rows × 4 columns

Improve dataset: Remove Gender column which is not affecting the result.

```
In [4]: data = data.drop(columns="Gender");
display(data);
```

	Height	Weight	Index
0	174	96	4
1	189	87	2
2	185	110	4
3	195	104	3
4	149	61	3
...
495	150	153	5
496	184	121	4
497	141	136	5
498	150	95	5
499	173	131	5

500 rows × 3 columns

Summary of each Features data:

```
In [5]: h = data["Height"].describe();  
w = data["Weight"].describe();  
i = data["Index"].describe();  
print(h);  
print("\n");  
print(w);  
print("\n");  
print(i);
```

```
count    500.000000
mean     169.944000
std       16.375261
min      140.000000
25%      156.000000
50%      170.500000
75%      184.000000
max      199.000000
Name: Height, dtype: float64
```

```
count    500.000000
mean     106.000000
std       32.382607
min       50.000000
25%       80.000000
50%      106.000000
75%      136.000000
max      160.000000
Name: Weight, dtype: float64
```

```
count    500.000000
mean       3.748000
std        1.355053
min         0.000000
25%         3.000000
50%         4.000000
75%         5.000000
max         5.000000
Name: Index, dtype: float64
```

Duplicat data in dataset:

```
In [6]: display(data[data.duplicated()]);
```

	Height	Weight	Index
20	157	110	5
162	192	101	3
187	182	84	3
197	177	117	4
260	159	104	5
310	171	147	5
321	181	111	4
327	167	85	4
334	157	56	2
347	162	58	2
354	190	50	0
355	174	90	3
365	141	80	5
381	191	62	1
382	177	117	4
395	164	71	3
398	149	61	3
400	195	104	3
419	177	61	2
421	140	146	5
462	179	56	1
466	188	99	3
482	142	86	5
492	198	50	0

Drop duplicate values:

```
In [7]: data.drop_duplicates(subset=None, inplace=True);
display(data[data.duplicated()]);
```

Height	Weight	Index
--------	--------	-------

There is no more duplicate values.

```
In [8]: h = data["Height"].describe();
w = data["Weight"].describe();
i = data["Index"].describe();
print(h);
```

```
print("\n");
print(w);
print("\n");
print(i);
```

```
count    476.000000
mean     169.878151
std       16.332011
min       140.000000
25%       156.000000
50%       170.000000
75%       184.000000
max       199.000000
Name: Height, dtype: float64
```

```
count    476.000000
mean     106.920168
std       32.319945
min       50.000000
25%       80.000000
50%       107.000000
75%       137.000000
max       160.000000
Name: Weight, dtype: float64
```


```
count    476.000000
mean       3.779412
std        1.337585
min         0.000000
25%         3.000000
50%         4.000000
75%         5.000000
max         5.000000
Name: Index, dtype: float64
```

Index Distribution

```
In [9]: plt.title("Index Distribution");
sns.countplot(data["Index"]);
```

C:\Users\mgone\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```




Height Distribution

```
In [10]: plt.figure(figsize=(30, 20));
plt.title("Height Distribution");
sns.countplot(data["Height"]);
```

C:\Users\mgone\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(




Weight Distribution

```
In [11]: plt.figure(figsize=(80, 20));
```

```
plt.title("Weight Distribution");
sns.countplot(data["Weight"]);
```

C:\Users\mgone\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```



Data on the 2D data space.


```
In [12]: x = data["Height"].values.tolist();
y = data["Weight"].values.tolist();

plt.figure(figsize=(50, 20));
plt.title("Height-Weight Graphic")
plt.scatter(x, y, label= "dot", color= "red",
            marker= ".", s=20)

plt.plot(x, y, "bo")

plt.xlabel("Height")
plt.ylabel("Weight")

plt.show()
```



Linear Regression Model

Training

We split height, weight and index values to equal two parts. We will use first half of the data to training our model. Data is already mixed. We can take first half of it as training data and the second half as test data.

```
In [13]: D = data.values.tolist();

X = data["Height"].values.tolist();
Y = data["Weight"].values.tolist();
Z = data["Index"].values.tolist();

x_train = X[:len(X)//2];
y_train = Y[:len(Y)//2];
z_train = Z[:len(Z)//2];

x_test = X[len(X)//2:];
y_test = Y[len(Y)//2:];
z_test = Z[len(Z)//2:];
```

Create Linear Regression Decider and calculate Loss, m1, m2 and b values using epoch and learning rate.

```
In [14]: lrm = LinearRegressionModel(learning_rate=0.000005, epoch=1000);
return_lists = lrm.fit(x_train, y_train, z_train, x_test, y_test, z_test);
loss_values = return_lists[0];
accuracy_values = return_lists[1];

y=np.arange(1,1001);

plt.figure(figsize=(50, 20));
plt.plot(y, loss_values);

plt.title("Change of Loss Values in Training");
plt.xlabel("Epoch Number");
plt.ylabel("Loss Values");

plt.scatter(y, loss_values, label= "dot", color= "red", marker= ".", s=25);

plt.show();

plt.figure(figsize=(50, 20));
plt.plot(y, accuracy_values);

plt.title("Change of Accuracy Values in Training");
plt.xlabel("Epoch Number");
plt.ylabel("Accuracy Values");


plt.scatter(y, accuracy_values, label= "dot", color= "red", marker= ".", s=25);

plt.show();

print("Description of Loss Values in Training");
df = pd.DataFrame(loss_values);
x = df.describe();
print(x);
print();
print("Description of Accuracy Values in Training");
df2 = pd.DataFrame(accuracy_values);
```



```
y = df2.describe();
print(y);
```



Description of Loss Values in Training

```
0
count    1000.000000
mean      135.167716
std       2620.847244
min        0.683892
25%        0.684526
50%        0.724378
75%        3.296848
max       76424.911765
```

Description of Accuracy Values in Training

```
0
count    1000.000000
mean       2.092425
std        7.040328
min        0.668142
25%        0.670595
50%        0.701266
75%        1.539085
max       170.266689
```

Testing

```

In [15]: loss_valuesT = return_lists[2];
accuracy_valuesT = return_lists[3];

y=np.arange(1,1001);

plt.figure(figsize=(50, 20));
plt.plot(y, loss_valuesT);

plt.title("Change of Loss Values in Testing");
plt.xlabel("Epoch Number");
plt.ylabel("Loss Values");

plt.scatter(y, loss_valuesT, label= "dot", color= "red", marker= ".", s=25);

plt.show();

plt.figure(figsize=(50, 20));
plt.plot(y, accuracy_valuesT);


plt.title("Change of Accuracy Values in Testing");
plt.xlabel("Epoch Number");
plt.ylabel("Accuracy Values");


plt.scatter(y, accuracy_valuesT, label= "dot", color= "red", marker= ".", s=25);

plt.show();

print("Description of Loss Values in Testing");
df = pd.DataFrame(loss_valuesT);
x = df.describe();
print(x);
print();
print("Description of Accuracy Values in Testing");
df2 = pd.DataFrame(accuracy_valuesT);
y = df2.describe();
print(y);

```





Description of Loss Values in Testing

```

0
count    1000.000000
mean      134.763438
std       2580.324732
min        0.555005
25%        0.556023
50%        0.603803
75%        3.550252
max       75193.000000

```

Description of Accuracy Values in Testing

```

0
count    1000.000000
mean       2.157303
std        7.084272
min        0.580831
25%        0.583156
50%        0.622557
75%        1.652211
max       169.131924

```

```

In [16]: prediction = lrm.predict(x_test, y_test);

prediction_result = [round(num) for num in prediction]
expected_result = z_test;

print("Test Features Expected Results")
print(expected_result)
print()
print("Prediction Results")
print(prediction_result)
print()
print("Actual Prediction Results")
print(prediction)

```

Test Features Expected Results

[1, 5, 0, 5, 3, 5, 3, 4, 3, 5, 5, 5, 5, 2, 4, 5, 5, 4, 5, 5, 2, 4, 5, 5, 5, 5, 1, 5, 5, 4, 0, 3, 3, 4, 2, 3, 1, 1, 5, 5, 4, 4, 4, 4, 5, 2, 5, 4, 3, 3, 4, 5, 5, 2, 4, 3, 4, 5, 4, 2, 4, 5, 4, 5, 5, 1, 5, 5, 5, 2, 2, 5, 3, 5, 4, 5, 4, 5, 5, 4, 2, 2, 3, 3, 5, 4, 2, 2, 2, 5, 5, 4, 5, 3, 4, 4, 3, 4, 4, 2, 5, 2, 2, 2, 2, 5, 4, 5, 1, 4, 1, 4, 5, 4, 5, 3, 4, 5, 4, 3, 5, 1, 2, 4, 5, 5, 5, 5, 3, 5, 5, 5, 2, 5, 4, 3, 2, 2, 2, 2, 3, 5, 3, 5, 5, 4, 2, 4, 4, 5, 2, 5, 5, 5, 1, 4, 5, 5, 5, 4, 5, 2, 5, 1, 5, 4, 1, 1, 4, 4, 4, 4, 2, 5, 5, 4, 2, 5, 5, 5, 1, 5, 4, 2, 5, 5, 4, 5, 4, 4, 5, 5, 5, 4, 5, 0, 2, 2, 4, 2, 4, 5, 4, 5, 5, 2, 5, 5, 5, 3, 5, 5, 2, 5, 3, 4, 5, 2, 2, 5, 5, 4, 4, 4, 4, 5, 2, 4, 5, 2, 5, 1, 5, 5, 4, 5, 5, 5]

Prediction Results

[2, 6, 1, 4, 3, 6, 3, 3, 2, 5, 5, 6, 5, 2, 4, 4, 5, 5, 5, 5, 2, 3, 3, 5, 5, 5, 2, 5, 4, 4, 2, 3, 3, 4, 3, 2, 2, 2, 5, 3, 4, 5, 5, 3, 6, 2, 4, 2, 4, 2, 3, 5, 6, 3, 3, 3, 3, 5, 4, 3, 4, 5, 4, 5, 3, 2, 5, 6, 5, 2, 2, 5, 3, 5, 3, 5, 4, 4, 5, 2, 2, 3, 2, 4, 6, 2, 2, 2, 2, 6, 4, 4, 5, 2, 4, 3, 3, 4, 4, 2, 5, 2, 3, 3, 2, 5, 4, 5, 2, 5, 2, 4, 6, 5, 5, 3, 4, 5, 2, 2, 4, 2, 2, 3, 5, 3, 5, 3, 3, 4, 6, 5, 2, 5, 3, 2, 2, 2, 2, 3, 3, 6, 3, 5, 5, 4, 2, 4, 5, 5, 2, 5, 4, 5, 2, 3, 6, 6, 4, 3, 4, 2, 5, 2, 5, 4, 2, 2, 5, 5, 3, 3, 3, 6, 3, 4, 2, 6, 5, 6, 2, 6, 4, 2, 5, 5, 5, 5, 2, 4, 4, 5, 6, 3, 4, 2, 3, 3, 5, 3, 3, 5, 4, 5, 5, 3, 5, 5, 4, 4, 5, 4, 2, 5, 3, 5, 5, 2, 2, 4, 4, 4, 3, 4, 4, 4, 3, 5, 5, 2, 5, 2, 3, 6, 4, 5, 3, 5]

Actual Prediction Results

[1.943214824741879, 5.574833703823112, 1.4977339083628736, 4.049631067999335, 3.3574428192177117, 5.812112521639207, 3.033741416969892, 3.0586890616329194, 2.342344343681884, 5.197709518196399, 5.051963129852098, 5.764969948706591, 5.251132301575292, 1.8567922403101547, 3.8066572224400677, 4.245462340625988, 5.206742445036115, 4.668352740988596, 4.67915761377184, 4.927826487074281, 1.6364089107672828, 2.898214718705729, 2.989747148177522, 4.965144957673566, 5.242099374735576, 5.325184060070758, 1.5330908380623354, 4.682891100615189, 3.621630822596657, 4.2356382382926565, 1.8157402828675195, 2.78644890186417, 3.1908778608005406, 3.983631465893673, 2.6420788717165875, 2.4910330433759205, 1.7568120667017495, 2.05458505383663, 4.52339752813791, 3.285163006665772, 3.7434103367278437, 4.920359513387581, 4.94628792850052, 2.5041950448057935, 5.765950719156503, 1.6387660394139136, 3.512617722148536, 2.490052272926009, 3.6259494921431106, 2.2437405282701977, 2.8170915742703704, 5.131709916090737, 5.653599719611839, 3.001137203663869, 2.60515598886411, 3.113488203208532, 2.8113965465271975, 4.97319711406337, 3.744391107177755, 2.6854879578058535, 3.828852150709656, 4.5021833703182335, 3.9624173080739955, 5.4833012743513185, 3.0600654198296384, 2.202292983080755, 5.047248872558836, 5.606457146679223, 5.193580443606242, 2.16497451248147, 1.8917535822628093, 4.73906660038752, 3.381994876133931, 4.883436630535104, 3.269643876589268, 5.092034316844821, 4.112482365964752, 3.657968522746031, 4.970839985416739, 2.2087791863175434, 1.7350127261789687, 3.0852026594489623, 2.3551107573649497, 3.505150748461837, 5.798950520209334, 2.406967587590827, 2.2366691423303053, 1.9880002690278642, 1.8345973120405659, 5.521015332697411, 4.483326341145187, 4.385703296183412, 5.225599474209161, 2.3763249151846266, 4.157267810250737, 3.1912734485473484, 2.7840917732175394, 4.043540452509355, 3.6622871922924847, 2.1753837975179047, 4.9146644856444075, 1.9459675411353174, 2.7723061299843854, 2.9351376015582065, 1.5906426960313862, 5.4809441457046875, 4.3470084673874085, 5.089281600451383, 1.8143639246708003, 5.282755744431404, 1.8998057386526133, 3.60041666477698, 5.752203535023525, 5.222846757815723, 5.442249316908683, 3.195592118093802, 4.470164339715313, 4.728657315351087, 2.2404026291736554, 2.037499970607111, 4.4153651981397015, 1.8799679390296553, 2.0479092556435456, 3.260215362002745, 5.212833060526096, 3.1808641635109134, 5.0694438008284255, 3.3460527637313655, 2.945942474341449, 4.2629430116023155, 5.597424219839508, 5.330879087813932, 2.3160203408221376, 4.990092602336593, 2.8288772175035244, 2.178136513911343, 1.5755191537016897, 1.8450065970770007, 2.2787018702228523, 2.5037994570589857, 2.929442573815033, 5.5950670911928775, 3.268267518392549, 4.858488985872076, 4.911326586547865, 3.7245533075547974, 2.273987612929591, 3.852423437175964, 4.569748925576911, 5.101462831431344, 2.1871694407510587, 5.13878130203063, 4.409670170396528, 4.749475885423956, 1.7393313957254226, 3.485312948838878, 5.5267103604405845, 5.5738529333732, 3.609845179363503, 3.388085491623912, 4.047273939352705, 2.0889612130861805, 4.642819913

622466, 1.8950914813593516, 5.205366086839396, 3.715124792968274, 1.673727381366568, 1.853058753466805, 5.007177685566113, 4.5838916974566954, 2.613208145253914, 3.281429 519822422, 2.6316695866801525, 5.569138676079938, 3.3941761071138927, 4.4866642402417 29, 2.0045001695542797, 5.712132348030802, 4.780118557830156, 5.635723460888705, 2.06 6370697069784, 5.523372461344042, 4.077521024012097, 2.390072099317604, 5.42280710503 2533, 5.285508460824843, 4.508859168511318, 5.036839587522402, 2.3588442442082993, 4. 261566653405596, 3.558383936884433, 5.230313731502423, 5.687580291114583, 3.071455475 315985, 4.278066553932011, 1.5000910370095042, 2.6198839434469985, 2.558994186381405 6, 4.504144911218057, 2.7935202878040624, 2.765630331791301, 5.248775172928661, 4.331 884925057712, 4.7437808576807825, 5.108534217371236, 2.663293029536265, 4.53616394182 0975, 5.147229046167241, 4.0647546103290315, 3.7453718776276665, 5.14251478887398, 3. 965170024467434, 2.089941983536092, 5.455015730591749, 2.657598001793091, 4.632015040 8392236, 5.102443601881256, 2.0771755698530265, 1.7666361690350805, 3.788195781013828 7, 4.078897382208816, 3.767962393644063, 2.939266676148364, 4.001112136870001, 3.8816 89751385445, 3.834942766199637, 2.6491502576564803, 4.790132255119783, 5.376645302549 829, 1.8907728118128977, 5.338346061500632, 1.6444610671570867, 3.4281566786166353, 5.561086519690135, 4.242709624232549, 4.924488587977738, 3.316786449521884, 4.6555863 27305532]

NOTE => The actual prediction values are used to calculate accuracy with Mean Error. The accuracy values are rounded to just show properly.

NOTE => In Training part, we have 1000 epoch. That's why, I did not print the output of the expected output and our output after each epoch turn.

Results

NOTE => Loss and Accuracy graphics are in the Trainig and Testing section of the report!

```
In [20]: asd = lrm.predict(x_test, y_test)
err = 0
for i in range (len(y_test)):
    err = err + abs(z_test[i] - asd[i])
err = err / len(y_test)
print("After we train our model, this is the last Mean Absolute Error Value for test c
print(err)
```

After we train our model, this is the last Mean Absolute Error Value for test data:
0.5808307740977067

m1, m2 and b Values

```
In [17]: r = lrm.values();
m1 = r[0]
m2 = r[1]
b = r[2]
print("m1: ", m1);
print("m2: ", m2);
print("b: ", b);
```

m1: -0.002357128646630796
m2: 0.03869482879600432
b: -0.005652989103906824

"Loss Values" and "Accuracy Values" charts for test and training data are available below the

relevant sections in the report.

When we start to train our linear regression model, during the first epochs accuracy values are like -170. This is because we randomly selected the initial assigned values for m_1 , m_2 , and b . After we train our model, accuracy values narrow down to 0.5 and 0.7 range. Because in every epoch, we calculate again loss value, m_1 , m_2 and b according to result of last epoch.

Same goes for loss values. During the first epochs it is too high. This is also because we randomly selected the initial assigned values for m_1 , m_2 and b . After the model is trained, accuracy values narrow down to 0.5 and 0.6 range as similar as accuracy values in training. But there is no overfitting. The accuracy value is lower than training as we expected. If we think about the overfitting graphic that we talked in our lectures, accuracy of training is usually more than testing.